



Sheaf-Theoretic Reward Spaces: System Architecture


Inputs



Reward Sheaf



Topological Analysis



Output



Riemannian Policy Gradient:

$$\nabla_{\theta} J = \mathbb{E}[\nabla \log \pi \cdot \frac{A}{g(s)}]$$

