

Constrained Policy Optimization

Summary

December 23, 2025

1 Problem Setup

Many real-world reinforcement learning problems involve safety constraints that must never be violated, such as robots avoiding dangerous states or autonomous systems maintaining safe operating conditions. Traditional RL methods cannot guarantee constraint satisfaction during learning. This paper introduces Constrained Policy Optimization (CPO), an algorithm that optimizes policies while ensuring that constraints are satisfied throughout the training process.

2 Thesis

The central thesis is that policy optimization can be performed safely by treating constraints as trust region problems and using primal-dual optimization to enforce them. The authors argue that by incorporating constraints directly into the policy update mechanism, agents can learn optimal behaviors while guaranteeing constraint satisfaction.

3 Key Concepts

Constrained RL: Reinforcement learning where the agent must maximize reward while satisfying one or more safety constraints that cannot be violated.

Trust Region Methods: Optimization approaches that limit the size of policy updates to ensure stable learning and constraint satisfaction.

Primal-Dual Optimization: A mathematical technique for solving constrained optimization problems by simultaneously optimizing the primal variables and dual variables (constraint multipliers).

Constraint Satisfaction: The requirement that certain conditions (safety constraints) must never be violated, not just on average but for all trajectories.

Policy Improvement: The process of updating a policy to improve performance while maintaining constraint satisfaction.

4 Methods

The authors develop Constrained Policy Optimization with several key components:

1. **Constrained Trust Region:** Policy updates are constrained to a trust region that ensures both performance improvement and constraint satisfaction.
2. **Primal-Dual Updates:** The algorithm uses primal-dual optimization to solve the constrained policy improvement problem analytically.

3. **Multiple Constraints:** The method can handle multiple simultaneous constraints, each with its own dual variable.
4. **Safe Exploration:** The trust region approach ensures that exploration does not violate constraints during learning.
5. **Theoretical Guarantees:** The algorithm provides theoretical guarantees for constraint satisfaction and monotonic improvement.

The key innovation is the integration of constraint enforcement directly into the policy optimization framework using trust region methods.

5 Results

The authors demonstrate CPO’s effectiveness on several constrained RL tasks:

- **Circle Navigation:** Agents learn to navigate around circular obstacles while staying within safe regions, achieving near-optimal performance while never violating constraints.
- **Gather Tasks:** In resource collection environments, agents learn to collect rewards while avoiding dangerous areas, demonstrating safe exploration and learning.
- **High-Dimensional Control:** The method scales to complex control tasks like Ant and Humanoid robots with multiple constraints.
- **Constraint Satisfaction:** Across all experiments, CPO maintains 100% constraint satisfaction while learning, unlike baseline methods that frequently violate constraints.
- **Performance:** CPO achieves comparable or better reward performance than unconstrained methods while satisfying constraints.
- **Sample Efficiency:** The method is sample efficient, requiring fewer episodes to learn safe policies compared to alternative approaches.

The results demonstrate that safe reinforcement learning is achievable without sacrificing performance.

6 Significance

This work represents a major advance in safe reinforcement learning:

- **Safety Guarantees:** Provides the first method with theoretical guarantees for constraint satisfaction during policy learning.
- **Practical Impact:** Makes RL applicable to safety-critical domains like robotics, autonomous vehicles, and healthcare where constraint violations are unacceptable.
- **Theoretical Foundation:** Establishes a rigorous mathematical framework for constrained policy optimization using trust region methods.
- **Scalability:** Demonstrates that safe RL methods can scale to high-dimensional control problems with multiple constraints.

- **Research Direction:** Opens up new research avenues for safe exploration, constraint learning, and real-world RL applications.

The paper provides a foundation for the development of safe and reliable reinforcement learning systems that can be deployed in real-world applications where safety is paramount.